


MEASURED EXECUTION

Every update. Less repeated work.

60 public projects, 180 commit-pinned release snapshots. Held-out evaluation: 20 projects, four workflow phases each. Every strategy sees the same source snapshots.

TEST WORKLOAD API COST (RECORDED USAGE + UNRESOLVED RESERVE)



Compare cost and acceptance together; quality-equivalent savings are not established.

reins

WHERE WORK IS SAVED

Four phases. One complete comparison.

Initial work is included. An unchanged recheck is shown separately, so cache savings are not confused with the cost of handling a real content change.

First collection

52 model calls

8 accepted | 13/20 assessed

Recorded Reins usage: \$0.08493. Parser/model/cache baseline: 27 calls, \$0.04592 recorded usage. Reins reused 0 stored field results.

Release update 1

45 model calls

7 accepted | 14/20 assessed

Recorded Reins usage: \$0.12793. Parser/model/cache baseline: 26 calls, \$0.09518 recorded usage. Reins reused 11 stored field results.

Release update 2

46 model calls

4 accepted | 12/20 assessed

Recorded Reins usage: \$0.13563. Parser/model/cache baseline: 26 calls, \$0.10936 recorded usage. Reins reused 13 stored field results.

Unchanged recheck

0 model calls

7 accepted | 13/20 assessed

Recorded Reins usage: \$0.00000. Parser/model/cache baseline: 0 calls, \$0.00000 recorded usage. Reins reused 64 stored field results.

Latency is observed processing time for frozen snapshots, including model calls, validation and persistence. Original public-source collection is excluded. Provider load and concurrent tasks can affect timing.

SOURCE-BACKED DELIVERABLES

A record, with its evidence.

Illustrative held-out outputs. Each value links to a pinned public source; the complete records are included in the raw experiment export.

celery/celery

PURPOSE

Task queues are used as a mechanism to distribute work across threads or machines.

INSTALL COMMAND

pip install -U Celery

LICENSE

BSD License

Open the pinned README

narwhals-dev/narwhals

PURPOSE

Extremely lightweight and extensible compatibility layer between dataframe libraries!

INSTALL COMMAND

pip install narwhals

LICENSE

MIT

Open the pinned README

reins

RECOVERY

Keep the response you already paid for.

Controlled fault injection using a real Gemini response, isolated from the release-change benchmark.

1. A model response and its usage were saved durably.
2. Execution was deliberately interrupted before the field checkpoint completed.
3. A new runtime resumed the task and reused that response.
4. Remaining work completed; subsequent unchanged phases reused saved results.

4 completed tasks

0 duplicate request keys

Project: joblib/joblib. Before restart: 1 call(s), \$0.00019. After completion: 5 total call(s), \$0.00299. Additional calls, if any, finish other fields; they do not repeat the interrupted request.

This validates one recovery boundary, not a production failure-rate reduction. Requests with uncertain provider outcomes remain held for reconciliation.

What the numbers mean.

The full experiment is saved locally. The public projection contains results and progress, not credentials, internal prompts or a paid execution endpoint.

QUALITY	Source support + separate Gemini Flash semantic judge; not human review. Reference annotations are model-assisted. The evaluator can make mistakes; no human-blind-review or statistical non-inferiority claim is made. Citation checks can reject formatting differences. 12/262 identical-answer groups received inconsistent model judgments.
SELECTION	20 development, 20 validation and 20 held-out projects; organization-isolated. The protocol requires 95% validation acceptance to select a baseline. No baseline met that requirement in this pilot. All strategies remain visible, including the zero-model parser.
COST	Input, cached-input and output token usage are priced explicitly. Failed or unresolved requests are retained; unresolved costs prevent a complete-cost savings claim. These are API estimates, not invoice debits. Local compute time is not converted into fictional cloud cost.
PRODUCT FOCUS	Incremental project updates, durable request reuse and task-level evidence. AWS AgentCore also supports observability, evaluation, configuration optimization and external models. This pilot is not a direct AWS performance benchmark.

references: \$1.37825 / 180 calls | evaluation: \$1.92799 / 240 calls | execution: \$4.97062 / 1379 calls
Recovery: \$0.00299. Recorded total: \$8.280; unresolved reserve: \$0.393. Labeling and evaluation overhead is reported separately.