

Incremental data refresh

Only changed fields are sent to Gemini; previously verified fields are reused.

72.0%

lower API cost per accepted result

300/300

Accepted test records

25.11s

p95 total time

Same task. Less model spend.

Before: one item per call

\$0.02971




After: changed fields only

\$0.00831



Estimated model cost per 1,000 accepted results

How it works



Also **69.1% lower cost than four-item batching**, with 300/300 accepted for both.

Selected examples from completed benchmarks, compared with Gemini 2.5 Flash Lite on the same 300 inputs. Constructed invoices and reformatted public metadata; p95 includes queueing. API list-price estimates, not billing debits.

Shared development/validation acquisition: \$0.06162, additional to execution costs. CPU and labeling are not priced. Structured inputs also support deterministic parsing. Results are measured examples, not guaranteed production savings.

Invoice extraction

Group independent invoice extractions into compact batches of 16.

7.7%

lower API cost per accepted result

300/300

Accepted test records

23.23s

p95 total time

Same task. Less model spend.

Before: one item per call


\$0.06432

After: batches of 16

\$0.05934

Estimated model cost per 1,000 accepted results

How it works



The batched configuration accepted 300/300 records, compared with 295/300 for one-item execution.

Selected examples from completed benchmarks, compared with Gemini 2.5 Flash Lite on the same 300 inputs. Constructed invoices and reformatted public metadata; p95 includes queueing. API list-price estimates, not billing debits.

Shared development/validation acquisition: \$0.08149, additional to execution costs. CPU and labeling are not priced. Structured inputs also support deterministic parsing. Results are measured examples, not guaranteed production savings.

Project research

Route ordinary records to Flash Lite, use Flash for conflicting sources, and repair invalid outputs.

49.6%

lower API cost per accepted result

300/300

Accepted test records

57.86s

p95 total time

Same task. Less model spend.

Before: one item per call

\$0.29247



After: model selection + batching

\$0.14737



Estimated model cost per 1,000 accepted results

How it works



300/300 accepted. p95 total: 163.39s before, 57.86s after. Observed p95 total time below 60 seconds.

Selected examples from completed benchmarks, compared with Gemini 2.5 Flash on the same 300 inputs. Constructed invoices and reformatted public metadata; p95 includes queueing. API list-price estimates, not billing debits.

Shared development/validation acquisition: \$0.09615, additional to execution costs. CPU and labeling are not priced. Structured inputs also support deterministic parsing. Results are measured examples, not guaranteed production savings.

Research paper collection

Batch public research metadata and select the model by source complexity. Best suited to offline collection.

63.5%

lower API cost per accepted result

300/300

Accepted test records

79.04s

p95 total time

Same task. Less model spend.

Before: one item per call

\$0.50033




After: model selection + batching

\$0.18271

Estimated model cost per 1,000 accepted results

How it works



300/300 accepted. p95 total: 239.91s before, 79.04s after. Offline batch workload; observed p95 total time is 79.04 seconds.

Selected examples from completed benchmarks, compared with Gemini 2.5 Flash on the same 300 inputs. Constructed invoices and reformatted public metadata; p95 includes queueing. API list-price estimates, not billing debits.

Shared development/validation acquisition: \$0.13314, additional to execution costs. CPU and labeling are not priced. Structured inputs also support deterministic parsing. Results are measured examples, not guaranteed production savings.